

Stability of the Lanczos Method

Christopher Musco (joint with Aaron Sidford and Cameron Musco)

Princeton University.

The Lanczos Method

The Lanczos Method

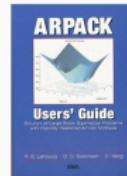
Used for solving linear systems, eigendecomposition, matrix exponentials, and approximating any matrix function.

The Lanczos Method

Used for solving linear systems, eigendecomposition, matrix exponentials, and approximating any matrix function.

- Introduced in 1950, developed through the 70s, ubiquitous in well-developed scientific computing libraries.

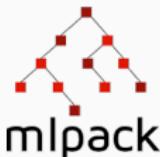
L A P A C K
L -A P -A C -K
L A P A -C -K
L -A P -A -C K
L A -P -A C K
L -A -P A C -K



The Lanczos Method

Used for solving linear systems, eigendecomposition, matrix exponentials, and approximating any matrix function.

- Introduced in 1950, developed through the 70s, ubiquitous in well-developed scientific computing libraries.
- Resurgence of interest due to new applications in data science and machine learning.



New applications combine Lanczos with super-scalable
stochastic iterative and randomized sketching methods.

New applications combine Lanczos with super-scalable
stochastic iterative and randomized sketching methods.

Require understanding performance under noisy inputs.

New applications combine Lanczos with super-scalable
stochastic iterative and randomized sketching methods.

Require understanding performance under noisy inputs.

Today's results:

1. Lanczos is very noise stable, performing essentially optimally for many matrix functions.

New applications combine Lanczos with super-scalable
stochastic iterative and randomized sketching methods.

Require understanding performance under noisy inputs.

Today's results:

1. Lanczos is very noise stable, performing essentially optimally for many matrix functions.
2. Except when solving linear systems! We provide a strong lower bound showing that noise can significantly impair Lanczos and the closely related conjugate gradient method.

STABILITY OF THE LANCZOS METHOD

The stability of Lanczos methods implemented in finite precision has been studied since the 70s and 80s.

STABILITY OF THE LANCZOS METHOD

The stability of Lanczos methods implemented in finite precision has been studied since the 70s and 80s.



Christopher
Paige



Anne
Greenbaum



Vladimir
Druskin

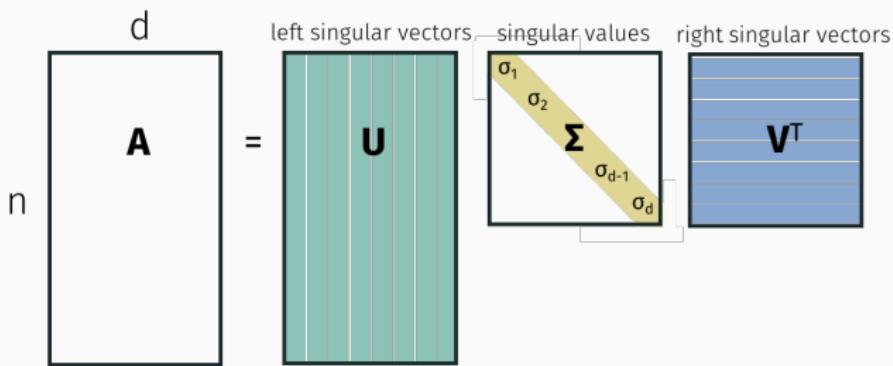


Leonid
Knizhnerman

WHAT IS A MATRIX FUNCTION?

WHAT IS A MATRIX FUNCTION?

Every matrix $A \in \mathbb{R}^{n \times d}$ has a singular value decomposition:



U, V are orthogonal, Σ is diagonal, $\sigma_1 \geq \dots \geq \sigma_d \in \mathbb{R}^+$.

WHAT IS A MATRIX FUNCTION?

Every symmetric matrix $A \in \mathbb{R}^{d \times d}$ has an orthogonal eigendecomposition:

$$d \begin{matrix} d \\ \boxed{\mathbf{A}} \end{matrix} = \begin{matrix} \text{eigenvectors} \\ \boxed{\mathbf{V}} \end{matrix} \begin{matrix} \text{eigenvalues} \\ \boxed{\Lambda} \end{matrix} \begin{matrix} \text{eigenvectors} \\ \boxed{\mathbf{V}^T} \end{matrix}$$

The diagram illustrates the orthogonal eigendecomposition of a symmetric matrix A . On the left, a square matrix A is shown with dimensions $d \times d$. To its right is an equals sign. Following the equals sign are three components: 1) A blue rectangle labeled "eigenvectors" containing the letter V . 2) A yellow diagonal rectangle labeled "eigenvalues" containing the letter Λ , with eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_{d-1}, \lambda_d$ written along the diagonal. 3) Another blue rectangle labeled "eigenvectors" containing the letter V^T .

WHAT IS A MATRIX FUNCTION?

For any scalar function $f : \mathbb{R} \rightarrow \mathbb{R}$ define $f(\mathbf{A})$:

$$d \begin{matrix} f(\mathbf{A}) \\ \boxed{} \end{matrix} = \begin{matrix} \text{eigenvectors} \\ \boxed{ \mathbf{V}} \end{matrix} \begin{matrix} \text{transformed} \\ \text{eigenvalues} \\ \boxed{ f(\Lambda)} \end{matrix} \begin{matrix} \text{eigenvectors} \\ \boxed{ \mathbf{V}^T} \end{matrix}$$

The diagram illustrates the computation of a matrix function $f(\mathbf{A})$. It shows the decomposition of the matrix \mathbf{A} into eigenvectors \mathbf{V} and eigenvalues Λ , followed by the application of the scalar function f to the eigenvalues. The resulting matrix $f(\mathbf{A})$ is then reconstructed by multiplying the transformed eigenvalues $f(\Lambda)$ by the transpose of the eigenvectors \mathbf{V}^T .

EXAMPLE MATRIX FUNCTION

$A^{-1} = f(A)$ where $f(\lambda) = 1/\lambda$

$$\overbrace{\left(\begin{bmatrix} v \\ v^T \end{bmatrix} \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix} \begin{bmatrix} v^T \\ v \end{bmatrix} \right)}^A \overbrace{\left(\begin{bmatrix} v \\ v^T \end{bmatrix} \begin{bmatrix} \frac{1}{\lambda_1} & & \\ & \ddots & \\ & & \frac{1}{\lambda_n} \end{bmatrix} \begin{bmatrix} v^T \\ v \end{bmatrix} \right)}^{A^{-1}} = I$$

COMPUTING MATRIX FUNCTIONS

Cost to compute $f(A)$:

COMPUTING MATRIX FUNCTIONS

Cost to compute $f(\mathbf{A})$:

$$\underbrace{O(n^3)}$$

eigendecompose $\mathbf{A} = \mathbf{V}\Lambda\mathbf{V}^T$

COMPUTING MATRIX FUNCTIONS

Cost to compute $f(\mathbf{A})$:

$$\underbrace{O(n^3)}_{\text{eigendecompose } \mathbf{A} = \mathbf{V}\Lambda\mathbf{V}^T} + \underbrace{O(n)}_{\text{compute } f(\Lambda)}$$

COMPUTING MATRIX FUNCTIONS

Cost to compute $f(\mathbf{A})$:

$$\underbrace{O(n^3)}_{\text{eigendecompose } \mathbf{A} = \mathbf{V}\boldsymbol{\Lambda}\mathbf{V}^T} + \underbrace{O(n)}_{\text{compute } f(\boldsymbol{\Lambda})} + \underbrace{O(n^3)}_{\text{form } \mathbf{V}f(\boldsymbol{\Lambda})\mathbf{V}^T}$$

COMPUTING MATRIX FUNCTIONS

Cost to compute $f(\mathbf{A})$:

$$\underbrace{O(n^3)}_{\text{eigendecompose } \mathbf{A} = \mathbf{V}\boldsymbol{\Lambda}\mathbf{V}^T} + \underbrace{O(n)}_{\text{compute } f(\boldsymbol{\Lambda})} + \underbrace{O(n^3)}_{\text{form } \mathbf{V}f(\boldsymbol{\Lambda})\mathbf{V}^T}$$

$$= O(n^3) \text{ in practice}$$

COMPUTING MATRIX FUNCTIONS

Cost to compute $f(\mathbf{A})$:

$$\underbrace{O(n^3)}_{\text{eigendecompose } \mathbf{A} = \mathbf{V}\Lambda\mathbf{V}^T} + \underbrace{O(n)}_{\text{compute } f(\Lambda)} + \underbrace{O(n^3)}_{\text{form } \mathbf{V}f(\Lambda)\mathbf{V}^T}$$

$$= O(n^3) \text{ in practice}$$

In theory can be improved to $O(n^\omega) \approx O(n^{2.3728639})$.
(but this is still slow)

FASTER MATRIX FUNCTIONS

Typically only interested in computing $f(\mathbf{A})\mathbf{x}$ for some $\mathbf{x} \in \mathbb{R}^n$.

$$f\left(\begin{bmatrix} & \\ & \mathbf{A} \\ & \end{bmatrix}\right) \cdot \begin{bmatrix} \mathbf{x} \end{bmatrix}$$

FASTER MATRIX FUNCTIONS

Typically only interested in computing $f(\mathbf{A})\mathbf{x}$ for some $\mathbf{x} \in \mathbb{R}^n$.

$$f\left(\begin{bmatrix} & \\ & \mathbf{A} \\ & \end{bmatrix}\right) \cdot \begin{bmatrix} \mathbf{x} \end{bmatrix}$$

Often much cheaper than computing $f(\mathbf{A})$ explicitly!
(this is what Lanczos and other algorithms target)

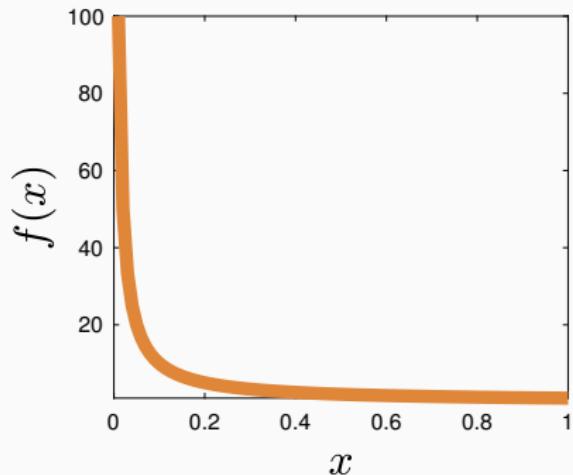
MATRIX INVERSE

Example

Linear system solving, $A^{-1}x$

Function

$$f(x) = 1/x$$



Countless applications...

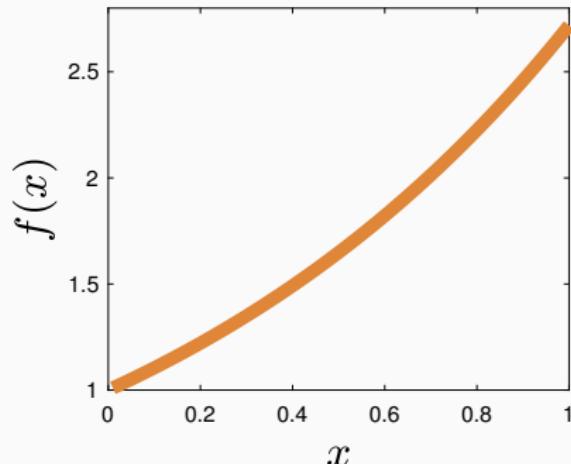
MATRIX EXPONENTIAL

Example

Matrix exponential, $e^A x$

Function

$$f(x) = e^x$$



Applications in semidefinite programming, graph algorithms
(balanced separator), differential equations.

[Arora, Hazan, Kale, '05], [Iyengar, Phillips, and Stein '11],
[Orecchia, Sachdeva, Vishnoi, '12], [Higham '08] (very complete survey)

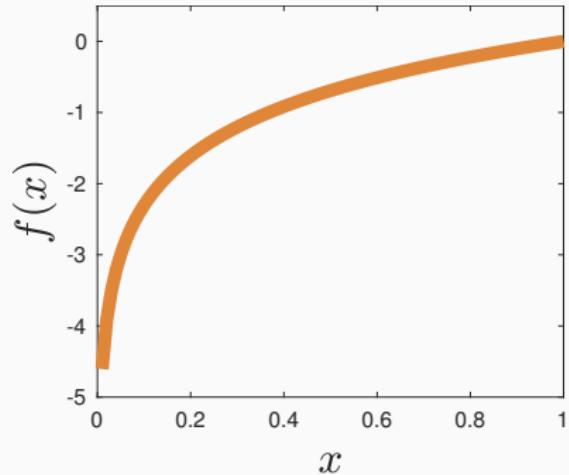
MATRIX LOG

Example

Matrix log, $\log(A)x$

Function

$$f(x) = \log(x)$$



Used to estimate $\log(\det(A)) = \text{tr}(\log(A))$.

Appears in log-likelihood equation for multivariate Gaussian.
Applications in Gaussian process regression, learning distance
kernels, Markov random fields.

[Dhillon, et al '06, '07, '08], [Han, Malioutov, Shin '15], [Saibaba, Alexanderian, Ipsen '17]

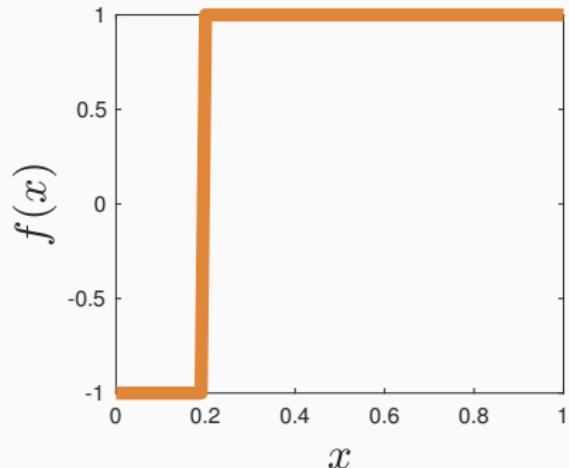
MATRIX STEP FUNCTION

Example

Step function, $\text{step}_\lambda(A)x$

Function

$$f(x) = \begin{cases} 1, & x \geq \lambda \\ 0, & x < \lambda \end{cases}$$



Projection to top eigenvectors, eigenvalue counting, computing matrix norms, spectral filtering, many more...

[Frostig, Musco, Musco, Sidford '16], [Saad, Ubaru '16], [Allen-Zhu, Li '17], [Tremblay, Puy, Gribonval, Vanderghenst '16], [Musco, Netrapalli, Sidford, Ubaru and Woodruff '18]

PRINCIPAL COMPONENT REGRESSION

Standard Regression:

Given: A, b

Solve: $x^* = \arg \min_x \|Ax - b\|^2$

PRINCIPAL COMPONENT REGRESSION

Standard Regression:

Given: A, b

Solve: $x^* = \arg \min_x \|Ax - b\|^2$

Principal Component Regression:

Given: A, b, λ

Solve: $x^* = \arg \min_x \|A_\lambda x - b\|^2$

PRINCIPAL COMPONENT REGRESSION

Standard Regression:

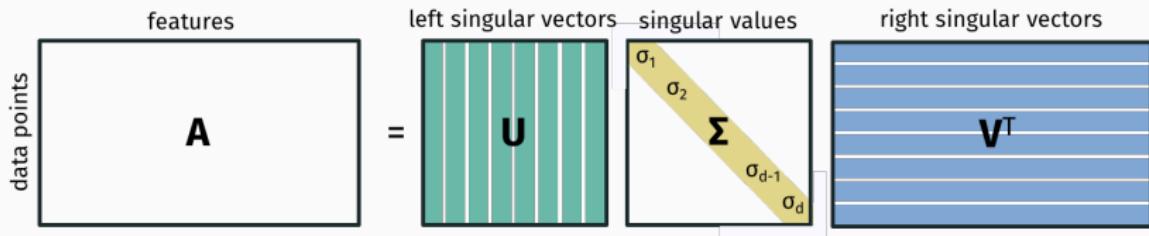
Given: A, b

Solve: $x^* = \arg \min_x \|Ax - b\|^2$

Principal Component Regression:

Given: A, b, λ

Solve: $x^* = \arg \min_x \|A_\lambda x - b\|^2$



PRINCIPAL COMPONENT REGRESSION

Standard Regression:

Given: A, b

Solve: $x^* = \arg \min_x \|Ax - b\|^2$

Principal Component Regression:

Given: A, b, λ

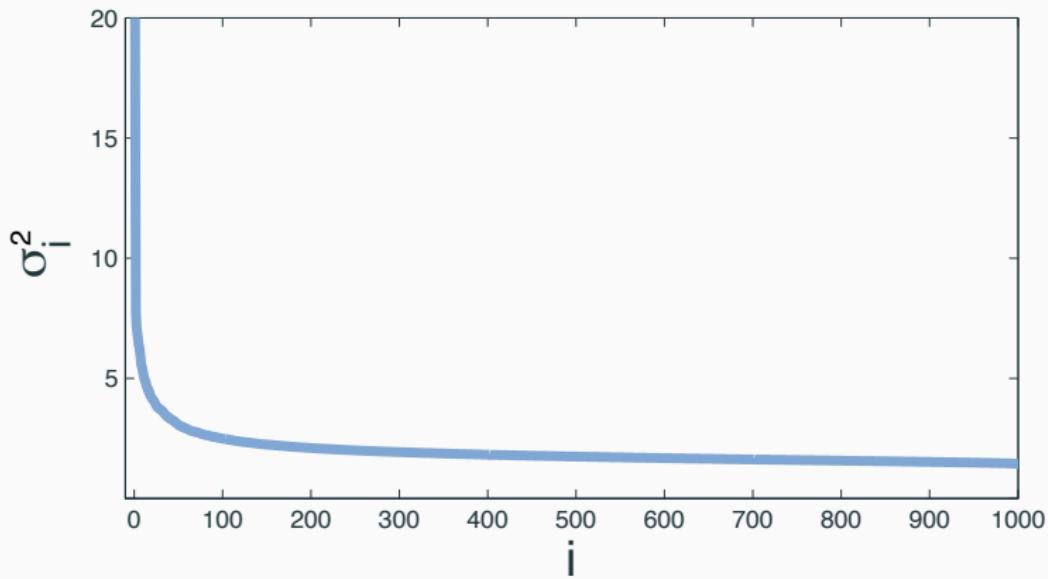
Solve: $x^* = \arg \min_x \|A_\lambda x - b\|^2$

$$\begin{matrix} & \text{features} \\ \text{data points} & \end{matrix} \quad \boxed{A_\lambda} \quad = \quad \begin{matrix} & \text{left singular vectors} \\ & \Sigma_\lambda \\ & \text{singular values} \end{matrix} \quad \begin{matrix} & \text{right singular vectors} \\ V_\lambda^\top \end{matrix}$$

The diagram illustrates the Singular Value Decomposition (SVD) of the matrix A_λ . On the left, a large rectangle labeled A_λ represents the data matrix. To its right is an equals sign. To the right of the equals sign is a vertical stack of three components: 1) 'left singular vectors' represented by a green vertical bar with three horizontal lines; 2) 'singular values' represented by a white square containing two yellow diagonal bars labeled σ_1 and σ_2 , with a smaller white square below it labeled Σ_λ ; 3) 'right singular vectors' represented by a blue vertical bar with four horizontal lines. The labels 'features' and 'data points' are positioned above and to the left of the matrix A_λ respectively.

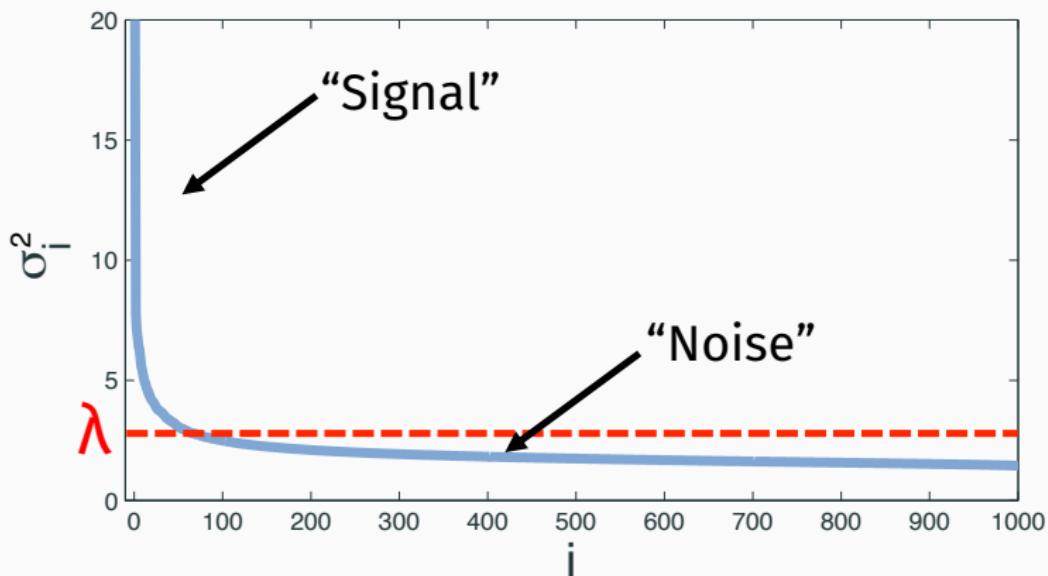
PRINCIPAL COMPONENT REGRESSION

Singular values of A



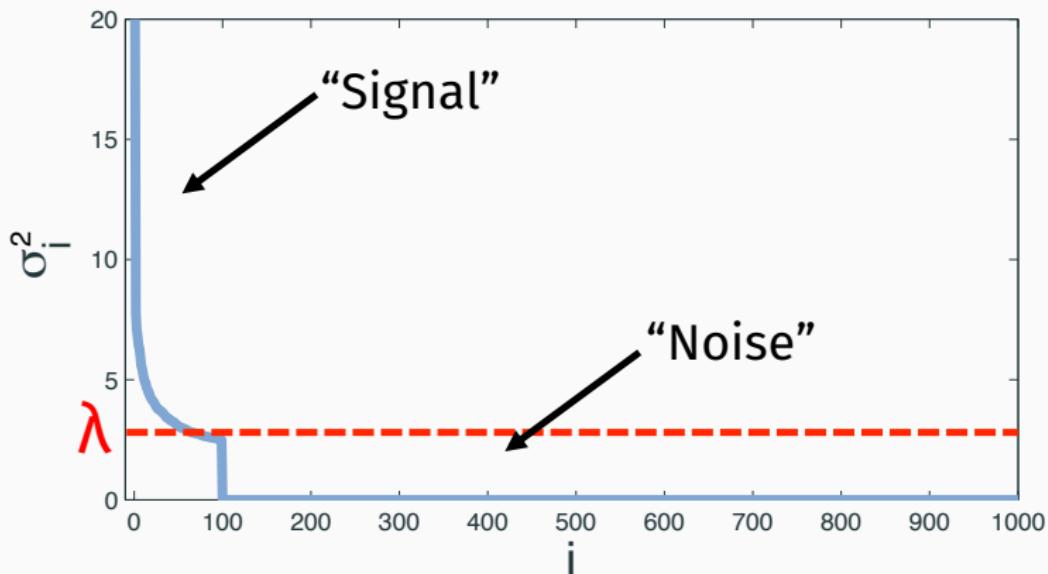
PRINCIPAL COMPONENT REGRESSION

Singular values of A



PRINCIPAL COMPONENT REGRESSION

Singular values of A_λ



Principal Component Regression (PCR):

$$\text{Goal: } \mathbf{x}^* = \arg \min_{\mathbf{x}} \|\mathbf{A}_{\lambda} \mathbf{x} - \mathbf{b}\|^2$$

$$\text{Solution: } \mathbf{x} = (\mathbf{A}_{\lambda}^T \mathbf{A}_{\lambda})^{-1} \mathbf{A}^T \mathbf{b}$$

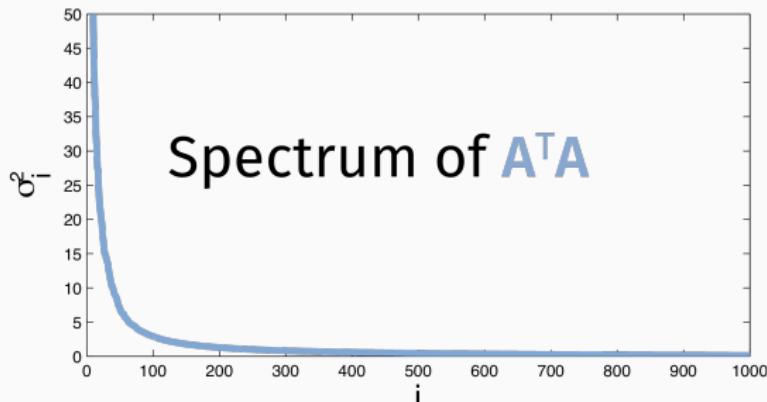
Principal Component Regression (PCR):

Goal: $\mathbf{x}^* = \arg \min_{\mathbf{x}} \|\mathbf{A}_\lambda \mathbf{x} - \mathbf{b}\|^2$

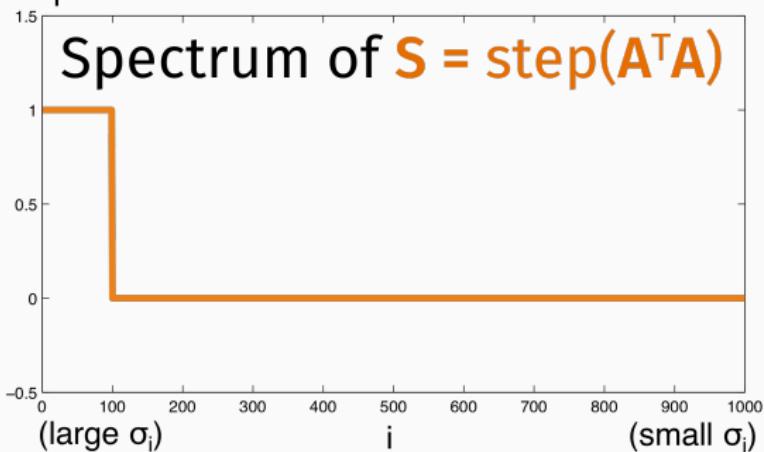
Solution: $\mathbf{x} = (\mathbf{A}_\lambda^T \mathbf{A}_\lambda)^{-1} \mathbf{A}^T \mathbf{b}$

Fastest way to apply $\mathbf{A}_\lambda^T \mathbf{A}_\lambda$ and $(\mathbf{A}_\lambda^T \mathbf{A}_\lambda)^{-1}$ to a vector is with a matrix step function.

PRINCIPAL COMPONENT REGRESSION



$$\mathbf{A}_\lambda^T \mathbf{A}_\lambda = \mathbf{S} \mathbf{A}^T \mathbf{A}$$



FAST ALGORITHMS FOR
MATRIX FUNCTIONS

MATRIX POLYNOMIALS

$$f \left(\begin{bmatrix} & A \\ & \end{bmatrix} \right) \cdot \begin{bmatrix} x \\ \end{bmatrix}$$

MATRIX POLYNOMIALS

Matrix **polynomials** can be computed iteratively.

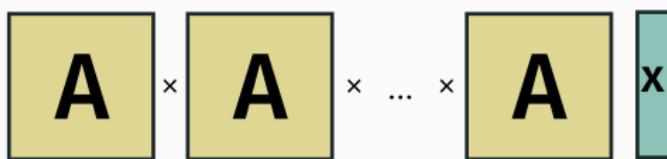
$$p \begin{pmatrix} & A \\ & \end{pmatrix} \cdot \begin{bmatrix} x \end{bmatrix}$$

MATRIX POLYNOMIALS

Matrix **polynomials** can be computed iteratively.

$$p \begin{pmatrix} & A \\ & \end{pmatrix} \cdot \begin{bmatrix} x \end{bmatrix}$$

$$A^k x = V \Lambda V^T V \Lambda V^T \cdots V \Lambda V^T x = V \Lambda^k V^T x$$

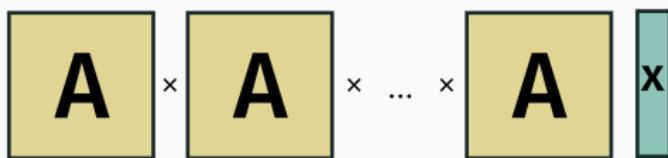


MATRIX POLYNOMIALS

Matrix **polynomials** can be computed iteratively.

$$p\left(\begin{bmatrix} & A \\ & \end{bmatrix}\right) \cdot \begin{bmatrix} x \end{bmatrix}$$

$$A^k x = V \Lambda V^T V \Lambda V^T \cdots V \Lambda V^T x = V \Lambda^k V^T x$$



Total time to compute $p(A)x = c_0x + c_1Ax + c_2A^2x + \dots + c_kA^kx$:

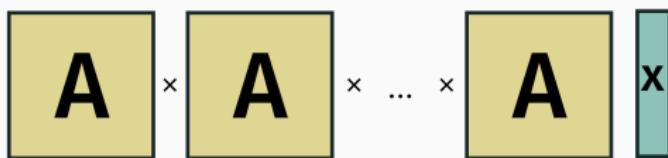
$$O(k \cdot \text{nnz}(A))$$

MATRIX POLYNOMIALS

Matrix **polynomials** can be computed iteratively.

$$p\left(\begin{bmatrix} & A \\ & \end{bmatrix}\right) \cdot \begin{bmatrix} x \end{bmatrix}$$

$$A^k x = V \Lambda V^T V \Lambda V^T \cdots V \Lambda V^T x = V \Lambda^k V^T x$$



Total time to compute $p(A)x = c_0x + c_1Ax + c_2A^2x + \dots + c_kA^kx$:

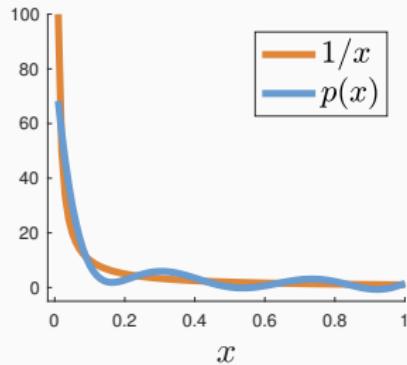
$$O(k \cdot \text{nnz}(A)) \leq O(k \cdot n^2) \ll O(n^3)$$

POLYNOMIAL APPROXIMATION

For general matrix functions:
approximate $f(x)$ with low-degree polynomial $p(x)$.

POLYNOMIAL APPROXIMATION

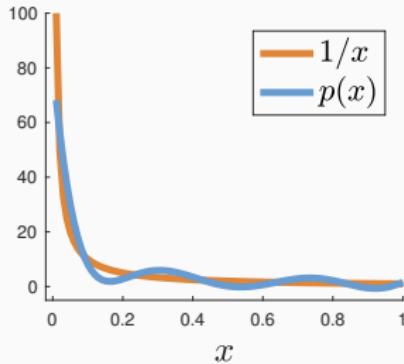
For general matrix functions:
approximate $f(x)$ with low-degree polynomial $p(x)$.



$$f(\mathbf{A})\mathbf{x} \approx p(\mathbf{A})\mathbf{x}$$

POLYNOMIAL APPROXIMATION

For general matrix functions:
approximate $f(x)$ with low-degree polynomial $p(x)$.



$$f(\mathbf{A})\mathbf{x} \approx p(\mathbf{A})\mathbf{x}$$

How does error in approximating scale function $f(\cdot)$ translate to error on matrix function?

POLYNOMIAL APPROXIMATION

$$\|f(A)x - p(A)x\| \leq \|f(A) - p(A)\| \cdot \|x\|$$

POLYNOMIAL APPROXIMATION

$$\|f(\mathbf{A})\mathbf{x} - p(\mathbf{A})\mathbf{x}\| \leq \|f(\mathbf{A}) - p(\mathbf{A})\| \cdot \|\mathbf{x}\| \leq \epsilon \cdot \|\mathbf{x}\|$$

where

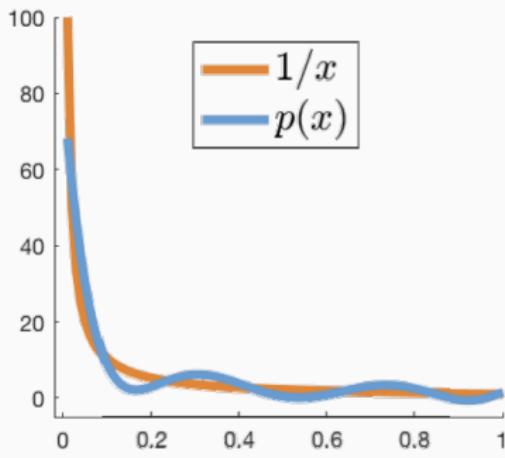
$$\epsilon = \max_{i=1,\dots,n} |f(\lambda_i) - p(\lambda_i)|.$$

POLYNOMIAL APPROXIMATION

$$\|f(A)x - p(A)x\| \leq \|f(A) - p(A)\| \cdot \|x\| \leq \epsilon \cdot \|x\|$$

where

$$\epsilon = \max_{i=1,\dots,n} |f(\lambda_i) - p(\lambda_i)|.$$

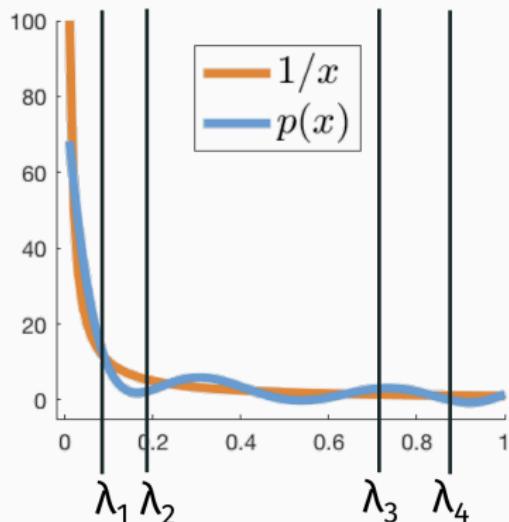


POLYNOMIAL APPROXIMATION

$$\|f(A)x - p(A)x\| \leq \|f(A) - p(A)\| \cdot \|x\| \leq \epsilon \cdot \|x\|$$

where

$$\epsilon = \max_{i=1,\dots,n} |f(\lambda_i) - p(\lambda_i)|.$$

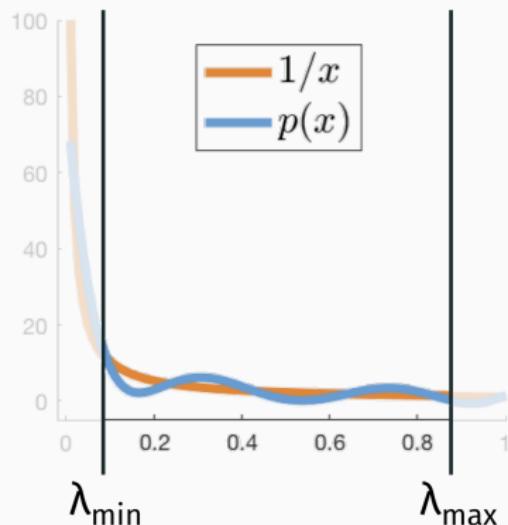


POLYNOMIAL APPROXIMATION

$$\|f(\mathbf{A})\mathbf{x} - p(\mathbf{A})\mathbf{x}\| \leq \|f(\mathbf{A}) - p(\mathbf{A})\| \cdot \|\mathbf{x}\| \leq \epsilon \cdot \|\mathbf{x}\|$$

where

$$\epsilon = \max_{i=1,\dots,n} |f(\lambda_i) - p(\lambda_i)|.$$



FINDING GOOD APPROXIMATING POLYNOMIALS

If we know $\lambda_{\min}(A)$ and $\lambda_{\max}(A)$ we can explicitly compute a near optimal polynomial p via Chebyshev interpolation.

FINDING GOOD APPROXIMATING POLYNOMIALS

If we know $\lambda_{\min}(A)$ and $\lambda_{\max}(A)$ we can explicitly compute a near optimal polynomial p via Chebyshev interpolation.

$$\delta_k = \min_{p \text{ a degree } k \text{ polynomial}} \left(\max_{x \in [\lambda_{\min}(A), \lambda_{\max}(A)]} |f(x) - p(x)| \right)$$

Final bound: Output y such that

$$\|f(A)x - y\| \leq O(\log k) \cdot \delta_k \cdot \|x\|.$$

FINDING GOOD APPROXIMATING POLYNOMIALS

If we know $\lambda_{\min}(A)$ and $\lambda_{\max}(A)$ we can explicitly compute a near optimal polynomial p via Chebyshev interpolation.

$$\delta_k = \min_{p \text{ a degree } k \text{ polynomial}} \left(\max_{x \in [\lambda_{\min}(A), \lambda_{\max}(A)]} |f(x) - p(x)| \right)$$

Final bound: Output y such that

$$\|f(A)x - y\| \leq O(\log k) \cdot \delta_k \cdot \|x\|.$$

Example bounds:

- Linear systems in $O\left(\sqrt{\lambda_{\max} / \lambda_{\min}}\right)$ iterations.
- Matrix exponential in $O(\|A\|)$ iterations.
- Matrix sign function in $O(1/\epsilon)$ iterations.
- Top eigenvector in $O(\log(n)/\sqrt{\epsilon})$ iterations.

Example bounds:

- Linear systems in $O\left(\sqrt{\lambda_{\max} / \lambda_{\min}}\right)$ iterations.
- Matrix exponential in $O(\|A\|)$ iterations.
- Matrix sign function in $O(1/\epsilon)$ iterations.
- Top eigenvector in $O(\log(n)/\sqrt{\epsilon})$ iterations.

But no one actually uses Chebyshev interpolation...

THE LANCZOS METHOD FOR MATRIX FUNCTIONS

THE LANCZOS METHOD



Cornelius Lanczos, 1950

THE LANCZOS METHOD



Cornelius Lanczos, 1950

- Simple to implement.
- No need to know $\lambda_{\min}(A)$ and $\lambda_{\max}(A)$.
- Much better convergence in practice (for many reasons).

THE LANCZOS METHOD



Cornelius Lanczos, 1950

- Simple to implement.
- No need to know $\lambda_{\min}(A)$ and $\lambda_{\max}(A)$.
- Much better convergence in practice (for many reasons).
- Matches optimal uniform approximation up to factor 2.

THE LANCZOS METHOD



Cornelius Lanczos, 1950

- Simple to implement.
- No need to know $\lambda_{\min}(A)$ and $\lambda_{\max}(A)$.
- Much better convergence in practice (for many reasons).
- Matches optimal uniform approximation up to factor 2.

Final bound: Output y such that $\|f(A)x - y\| \leq 2\delta_k \cdot \|x\|$.

LANCZOS METHOD FOR MATRIX FUNCTIONS

Form orthogonal span for Krylov subspace $\mathcal{K} = \{\mathbf{x}, \mathbf{A}\mathbf{x}, \dots, \mathbf{A}^k\mathbf{x}\}$:

$$Q = \text{span} \left(x, Ax, \dots, A^k x \right)$$

LANCZOS METHOD FOR MATRIX FUNCTIONS

Form orthogonal span for Krylov subspace $\mathcal{K} = \{\mathbf{x}, \mathbf{Ax}, \dots, \mathbf{A}^k \mathbf{x}\}$:

$$Q = \text{span} \left\{ x, Ax, \dots, A^k x \right\}$$

Runtime: $O(k \cdot \text{nnz}(\mathbf{A})) + O(nk^2)$

LANCZOS METHOD FOR MATRIX FUNCTIONS

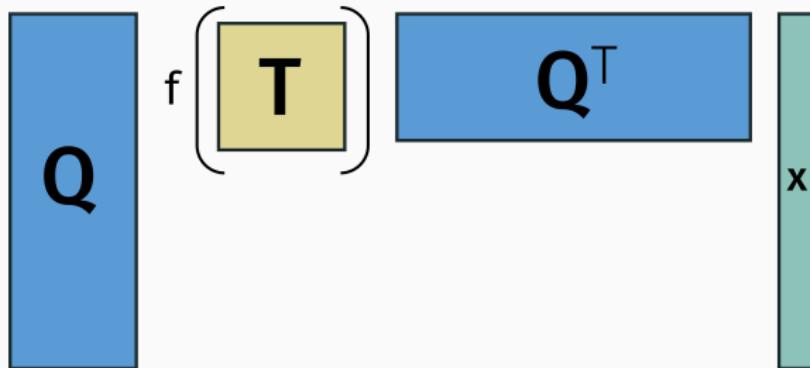
Compute:

$$\leftarrow \begin{matrix} k \\ \textcolor{brown}{T} \end{matrix} = \begin{matrix} \textcolor{blue}{Q}^T \\ \textcolor{teal}{A} \\ n \end{matrix} \quad \begin{matrix} n \\ \textcolor{blue}{Q} \end{matrix}$$

Runtime: $O(k \cdot \text{nnz}(A)) + O(nk^2)$

LANCZOS METHOD FOR MATRIX FUNCTIONS

Approximate:

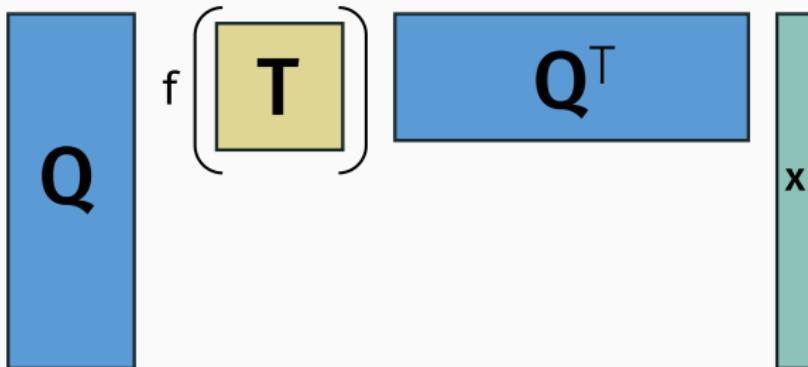


Runtime: $O(k \cdot \text{nnz}(A)) + O(nk^2) + O(k^3)$

Reduce the problem to the cost of computing a matrix function for a $k \times k$ matrix.

LANCZOS METHOD FOR MATRIX FUNCTIONS

Approximate:



~~Runtime: $O(k \cdot \text{nnz}(A)) + O(nk^2) + O(k^3)$~~

Runtime: $O(k \cdot \text{nnz}(A) + nk)$

Reduce the problem to the cost of computing a matrix function for a $k \times k$ matrix.

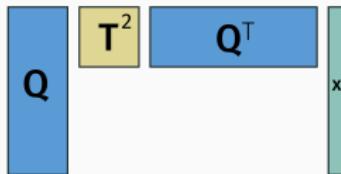
QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

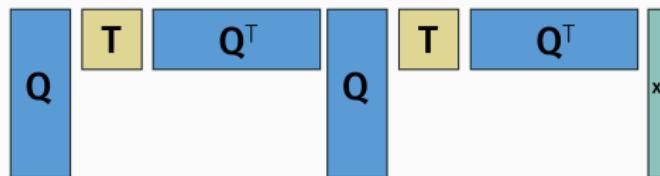
Proof:



QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

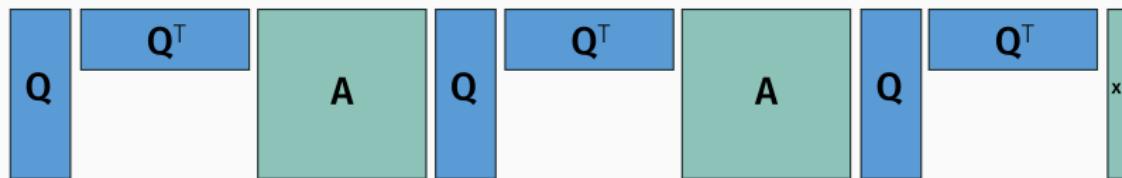
Proof:



QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

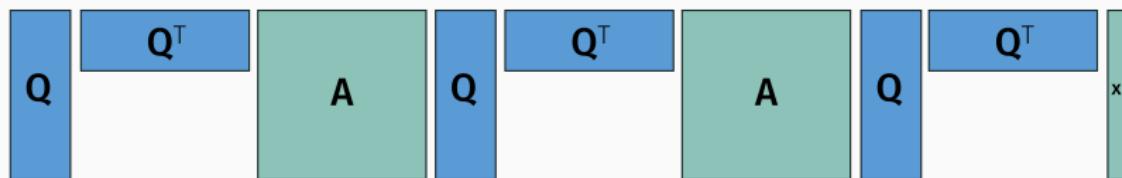
Proof:



QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

Proof:

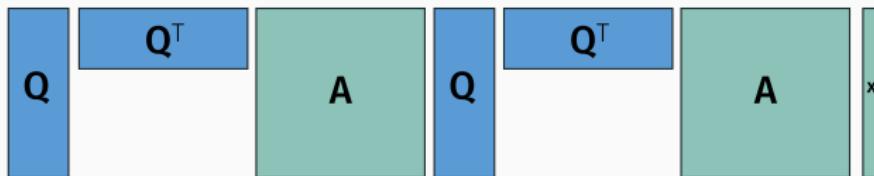


x, Ax, A^2x all lie in the span of Q .

QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

Proof:

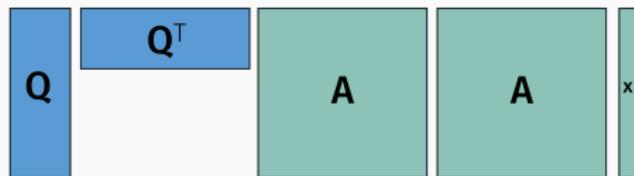


x, Ax, A^2x all lie in the span of Q .

QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

Proof:



x, Ax, A^2x all lie in the span of Q .

QUICK ANALYSIS OF LANCZOS

Claim: Lanczos applies degree k polynomials exactly.

Proof:



x, Ax, A^2x all lie in the span of Q .

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

Lanczos automatically applies the polynomial “part” of f .
(simple application of triangle inequality)

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

Lanczos automatically applies the polynomial “part” of f .
(simple application of triangle inequality)

For any degree k polynomial p ,

$$\begin{aligned}\|f(A)x - Qf(T)Q^T x\| &\leq \|f(A)x - p(A)x\| \\ &+ \|p(A)x - Qp(T)Q^T x\| \\ &+ \|Qp(T)Q^T x - Qf(T)Q^T x\|\end{aligned}$$

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

Lanczos automatically applies the polynomial “part” of f .
(simple application of triangle inequality)

For any degree k polynomial p ,

$$\begin{aligned}\|f(A)x - Qf(T)Q^T x\| &\leq \|f(A)x - p(A)x\| \\ &\quad + \|p(A)x - Qp(T)Q^T x\| \\ &\quad + \|Qp(T)Q^T x - Qf(T)Q^T x\| \\ &\leq \delta_k \|x\|\end{aligned}$$

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

Lanczos automatically applies the polynomial “part” of f .
(simple application of triangle inequality)

For any degree k polynomial p ,

$$\begin{aligned}\|f(A)x - Qf(T)Q^T x\| &\leq \|f(A)x - p(A)x\| \\ &\quad + \|p(A)x - Qp(T)Q^T x\| \\ &\quad + \|Qp(T)Q^T x - Qf(T)Q^T x\| \\ &\leq \delta_k \|x\| + 0\end{aligned}$$

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

Lanczos automatically applies the polynomial “part” of f .
(simple application of triangle inequality)

For any degree k polynomial p ,

$$\begin{aligned}\|f(A)x - Qf(T)Q^T x\| &\leq \|f(A)x - p(A)x\| \\ &\quad + \|p(A)x - Qp(T)Q^T x\| \\ &\quad + \|Qp(T)Q^T x - Qf(T)Q^T x\| \\ &\leq \delta_k \|x\| + 0 + \|p(T) - f(T)\| \cdot \|Q^T x\|\end{aligned}$$

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

Lanczos automatically applies the polynomial “part” of f .
(simple application of triangle inequality)

For any degree k polynomial p ,

$$\begin{aligned}\|f(\mathbf{A})\mathbf{x} - \mathbf{Q}f(\mathbf{T})\mathbf{Q}^T\mathbf{x}\| &\leq \|f(\mathbf{A})\mathbf{x} - p(\mathbf{A})\mathbf{x}\| \\ &\quad + \|p(\mathbf{A})\mathbf{x} - \mathbf{Q}p(\mathbf{T})\mathbf{Q}^T\mathbf{x}\| \\ &\quad + \|\mathbf{Q}p(\mathbf{T})\mathbf{Q}^T\mathbf{x} - \mathbf{Q}f(\mathbf{T})\mathbf{Q}^T\mathbf{x}\| \\ &\leq \delta_k \|\mathbf{x}\| + 0 + \delta_k \|\mathbf{x}\|.\end{aligned}$$

Since $\mathbf{T} = \mathbf{Q}^T \mathbf{A} \mathbf{Q}$, $[\lambda_{\min}(\mathbf{T}), \lambda_{\max}(\mathbf{T})] \subseteq [\lambda_{\min}(\mathbf{A}), \lambda_{\max}(\mathbf{A})]$.

QUICK ANALYSIS OF LANCZOS

How about for a general functions $f(x)$?

Lanczos automatically applies the polynomial “part” of f .
(simple application of triangle inequality)

For any degree k polynomial p ,

$$\begin{aligned}\|f(A)x - Qf(T)Q^T x\| &\leq \|f(A)x - p(A)x\| \\ &\quad + \|p(A)x - Qp(T)Q^T x\| \\ &\quad + \|Qp(T)Q^T x - Qf(T)Q^T x\| \\ &\leq \delta_k \|x\| + 0 + \delta_k \|x\|. \\ &= 2\delta_k \cdot \|x\|\end{aligned}$$

POLYNOMIAL METHODS WITH NOISE

MATRIX FUNCTIONS WITH NOISE

In many applications, we do not multiply by A exactly!

MATRIX FUNCTIONS WITH NOISE

In many applications, we do not multiply by A exactly!

$$\boxed{A} \quad \boxed{x} + \boxed{\epsilon_1}$$

MATRIX FUNCTIONS WITH NOISE

In many applications, we do not multiply by A exactly!

$$\mathbf{A} \left(\mathbf{A} \mathbf{x} + \boldsymbol{\epsilon}_1 \right) + \boldsymbol{\epsilon}_2$$

MATRIX FUNCTIONS WITH NOISE

In many applications, we do not multiply by A exactly!

$$A \begin{pmatrix} A \\ A \end{pmatrix} \left(A \begin{pmatrix} x \\ \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \end{pmatrix} + \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \end{pmatrix} \right)$$

MATRIX FUNCTIONS WITH NOISE

In many applications, we do not multiply by A exactly!

$$A \underbrace{A}_{\text{A}} \left(A \underbrace{x + \epsilon_1}_{\text{x} + \epsilon_1} + \epsilon_2 \right) + \epsilon_3$$

Natural model when Lanczos is combined with scalable
randomized or iterative methods.

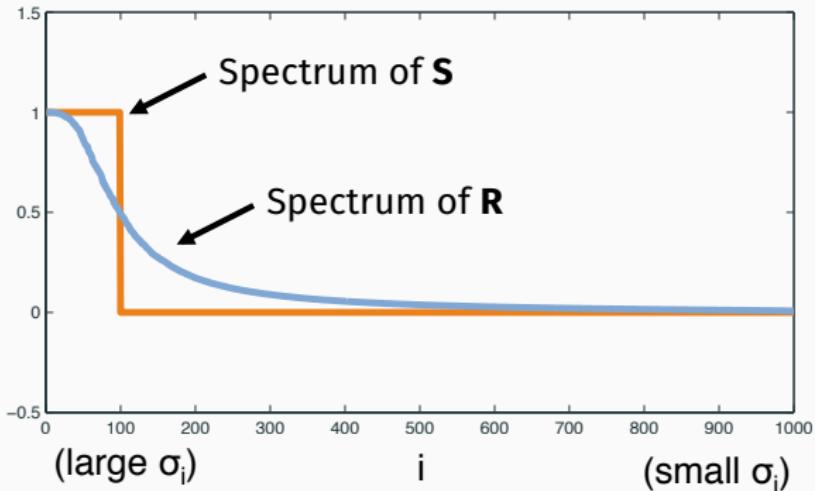
MATRIX FUNCTIONS WITH NOISE

Powerful paradigm:

- $A = B^{-1}$ for some matrix B .
- Apply B^{-1} to vectors quickly and approximately.

MATRIX STEP FUNCTION

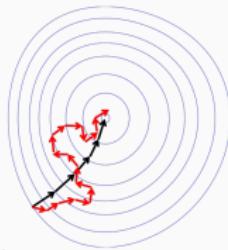
Fastest algorithms for computing $S = \text{step}_\lambda(A^T A)$ actually compute $\text{step}_{1/2}(R)$ where $R = (A^T A + \lambda I)^{-1} A^T A$.



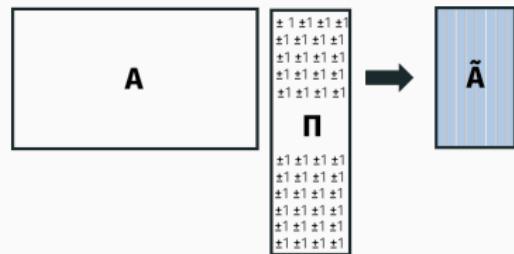
Most of the work is computing Rx .

LANCZOS AND FAST MATRIX INVERSION

$(\mathbf{A}^T \mathbf{A} + \lambda \mathbf{I})^{-1}$ can be computed approximately with a variety of highly (often randomized) scalable methods:



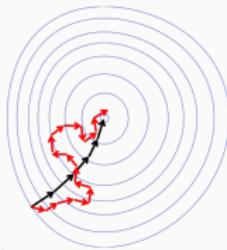
Stochastic Iterative
Methods



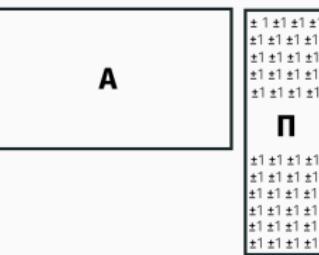
Randomized Sketching

LANCZOS AND FAST MATRIX INVERSION

$(\mathbf{A}^T \mathbf{A} + \lambda \mathbf{I})^{-1}$ can be computed approximately with a variety of highly (often randomized) scalable methods:



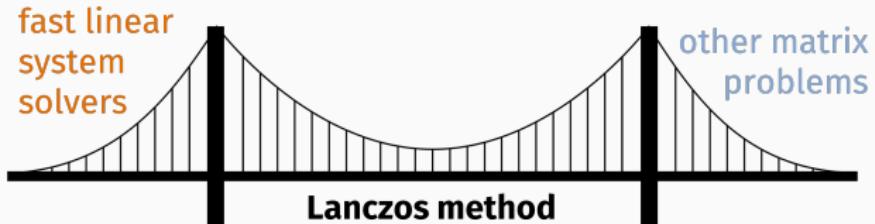
Stochastic Iterative
Methods



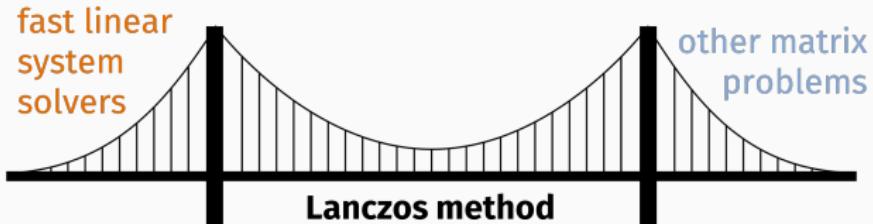
Randomized Sketching

Runtimes scale roughly as $O(\text{nnz}(\mathbf{A}) \cdot \log(1/\epsilon))$.
(for ϵ approximate solution)

LANCZOS AND FAST MATRIX INVERSION



LANCZOS AND FAST MATRIX INVERSION



- Faster eigenvector algorithms (in many regimes).
- Faster eigenvalue counting algorithms.
- Faster log-determinant and matrix norm algorithms.
- Faster balanced separator algorithms for graphs (via Laplacian matrix exponential).

Are matrix function algorithms stable?

We need to understand how the performance of our algorithms change when we replace every matrix-vector multiplication \mathbf{Ax} with an approximate solution.

Are matrix function algorithms stable?

We need to understand how the performance of our algorithms change when we replace every matrix-vector multiplication \mathbf{Ax} with an approximate solution.

Same stability questions were asked decades ago to understand roundoff error when computing \mathbf{Ax} .

$$fl(x \circ y) = (1 \pm \epsilon)(x \circ y) \text{ for } \circ = +, -, \times, \div$$

It is very easy to design iterative methods that converge very slowly when \mathbf{Ax} is computed approximately. But the Lanczos method (with no modifications) continues to perform well.

Can we explain this phenomena?

STABLE POLYNOMIAL COMPUTATION

Can we apply scalar polynomials in a stable way?

Can we apply scalar polynomials in a stable way?

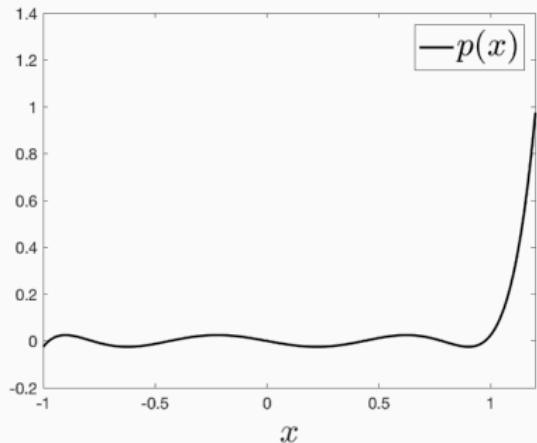
1. Want to compute $p(x) = c_0 + c_1x + \dots + c_kx^k$.
2. We do not know x , but we have access to a function `approxMult` that for any input z outputs:

$$\text{approxMult}(x, z) = z \cdot x + \epsilon.$$

STABLE POLYNOMIAL COMPUTATION

Goal: Compute $p(x) = 64x^7 - 112x^5 + 56x^3 - 7x$.

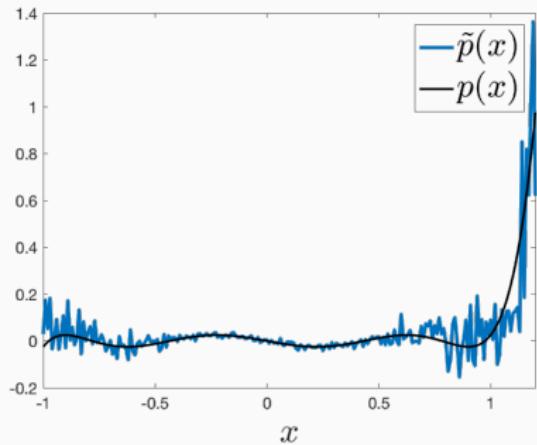
Using **approxMult** with $\epsilon = .05$.



STABLE POLYNOMIAL COMPUTATION

Goal: Compute $p(x) = 64x^7 - 112x^5 + 56x^3 - 7x$.

Using `approxMult` with $\epsilon = .05$.



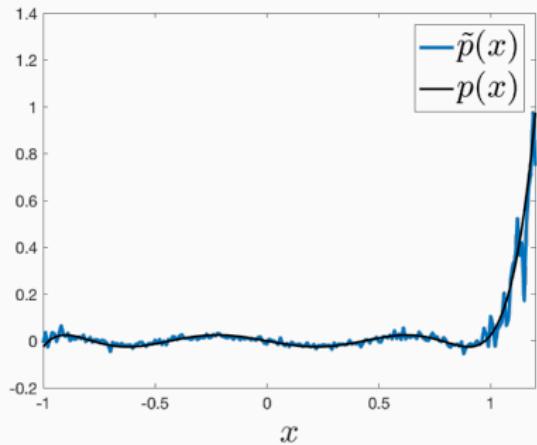
Directly compute and sum monomials.

$$x^i = \text{approxMult}(\text{approxMult}(\dots \text{approxMult}(1)\dots))$$

STABLE POLYNOMIAL COMPUTATION

Goal: Compute $p(x) = 64x^7 - 112x^5 + 56x^3 - 7x$.

Using `approxMult` with $\epsilon = .05$.



Factor $p(x) = (x - .98)(x - .78)\dots(x - .43)$.

$$t_1 = \text{approxMult}(1) - .98$$

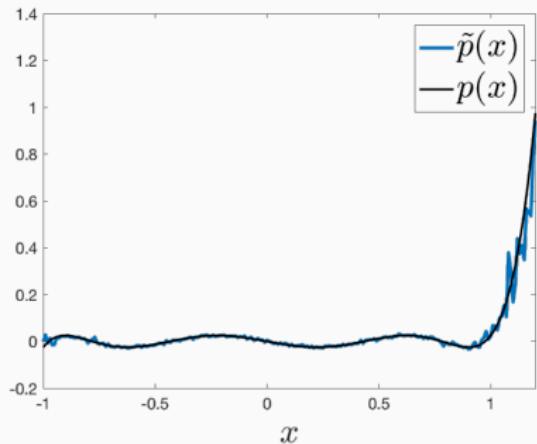
$$t_2 = \text{approxMult}(t_1) - .78 \cdot t_1$$

...

STABLE POLYNOMIAL COMPUTATION

Goal: Compute $p(x) = 64x^7 - 112x^5 + 56x^3 - 7x$.

Using `approxMult` with $\epsilon = .05$.



It's possible to do even better.

STABLE POLYNOMIAL COMPUTATION

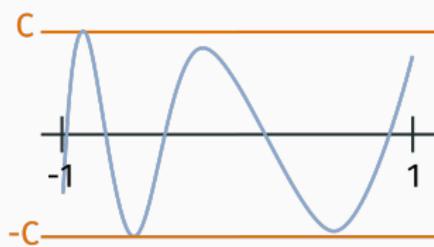
Assume we want to approximate $p(x)$ for $x \in [-1, 1]$.

Assume $|p(x)| \leq C$.

STABLE POLYNOMIAL COMPUTATION

Assume we want to approximate $p(x)$ for $x \in [-1, 1]$.

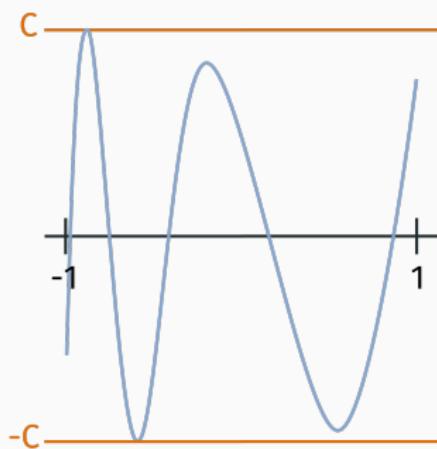
Assume $|p(x)| \leq C$.



STABLE POLYNOMIAL COMPUTATION

Assume we want to approximate $p(x)$ for $x \in [-1, 1]$.

Assume $|p(x)| \leq C$.



Claim

We can compute any $p(x)$ to accuracy $\epsilon \cdot Ck^3$ if **approxMult** has accuracy ϵ .

Claim

We can compute any $p(x)$ to accuracy $\epsilon \cdot Ck^3$ if `approxMult` has accuracy ϵ .

When translating to matrices:

- Runtime of linear system solvers scales with $\log(1/\epsilon)$

Claim

We can compute any $p(x)$ to accuracy $\epsilon \cdot Ck^3$ if `approxMult` has accuracy ϵ .

When translating to matrices:

- Runtime of linear system solvers scales with $\log(1/\epsilon)$
- Set $\epsilon = \epsilon'/Ck^3$ to get error ϵ' .

Claim

We can compute any $p(x)$ to accuracy $\epsilon \cdot Ck^3$ if `approxMult` has accuracy ϵ .

When translating to matrices:

- Runtime of linear system solvers scales with $\log(1/\epsilon)$
- Set $\epsilon = \epsilon'/Ck^3$ to get error ϵ' .
- Runtime overhead is just $O(\log(kC))$.

FIRST ATTEMPT

Compute monomials:

$$(x + \epsilon_1)$$

FIRST ATTEMPT

Compute monomials:

$$(x(x + \epsilon_1) + \epsilon_2)$$

FIRST ATTEMPT

Compute monomials:

$$(x(x(x + \epsilon_1) + \epsilon_2) + \epsilon_3)$$

FIRST ATTEMPT

Compute monomials:

$$x^i + x^{i-1}\epsilon_1 + x^{i-2}\epsilon_2 + \dots + \epsilon_j.$$

FIRST ATTEMPT

Compute monomials:

$$x^i + x^{i-1}\epsilon_1 + x^{i-2}\epsilon_2 + \dots + \epsilon_j.$$

Since $|x| \leq 1$, error on x^i bounded by $\epsilon_1 + \epsilon_2 + \dots + \epsilon_3 \leq \textcolor{brown}{\epsilon i}$.

FIRST ATTEMPT

Compute monomials:

$$x^i + x^{i-1}\epsilon_1 + x^{i-2}\epsilon_2 + \dots + \epsilon_j.$$

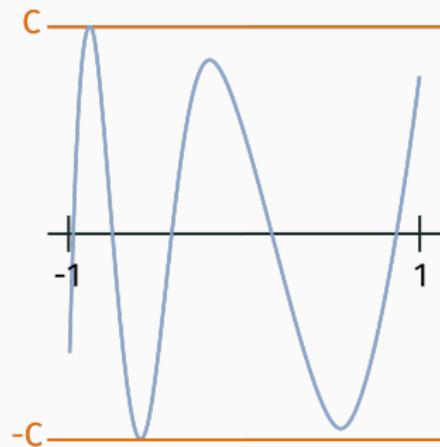
Since $|x| \leq 1$, error on x^i bounded by $\epsilon_1 + \epsilon_2 + \dots + \epsilon_j \leq \epsilon j$.

We can then compute $p(x) = c_0 + c_1x + \dots + c_kx^k$ up to error:

$$c_1\epsilon + 2 \cdot c_2\epsilon + \dots + k \cdot c_k\epsilon \leq \epsilon k \cdot \sum_{i=1}^k |c_i|$$

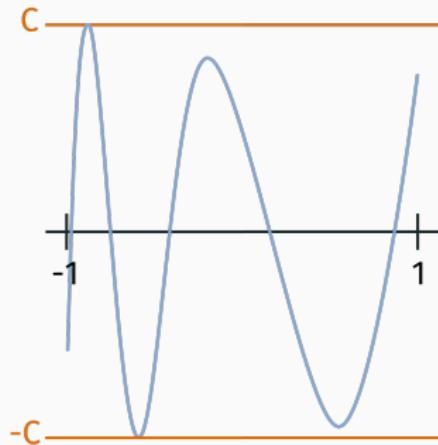
FIRST ATTEMPT

$\sum_{i=1}^k |c_k|$ can be far larger than our goal of $\epsilon \cdot Ck^3$.



FIRST ATTEMPT

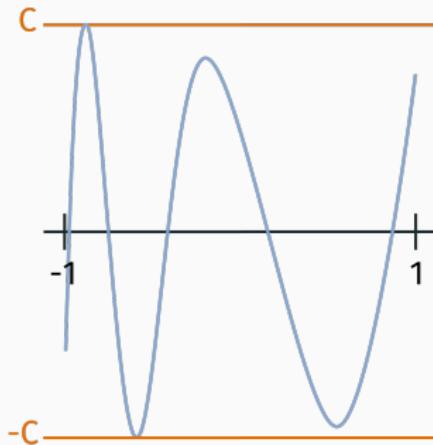
$\sum_{i=1}^k |c_k|$ can be far larger than our goal of $\epsilon \cdot Ck^3$.



There are polynomials with $C = 1$ but $\sum_{i=1}^k |c_i| = O(2^k)$.

FIRST ATTEMPT

$\sum_{i=1}^k |c_k|$ can be far larger than our goal of $\epsilon \cdot Ck^3$.



There are polynomials with $C = 1$ but $\sum_{i=1}^k |c_k| = O(2^k)$.

Exponential instead of polynomial loss in k .

$$\log\left(\frac{1}{\epsilon/2^k}\right) = k + \log(1/\epsilon)$$

“BAD” POLYNOMIALS

What are those polynomials?

“BAD” POLYNOMIALS

What are those polynomials?

Chebyshev polynomials of the first kind.

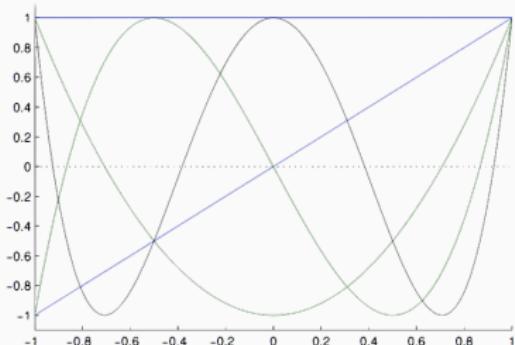
$$T_0(x) = 1$$

$$T_1(x) = x$$

$$T_2(x) = 2x^2 - 1$$

⋮

$$T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$$



“BAD” POLYNOMIALS

What are those polynomials?

Chebyshev polynomials of the first kind.

$$T_0(x) = 1$$

$$T_1(x) = x$$

$$T_2(x) = 2x^2 - 1$$

$$T_3(x) = 4x^3 - 3x$$

$$T_4(x) = 8x^4 - 8x^2 + 1$$

$$T_5(x) = 16x^5 - 20x^3 + 5x$$

$$T_6(x) = 32x^6 - 48x^4 + 18x^2 - 1$$

$$T_7(x) = 64x^7 - 112x^5 + 56x^3 - 7x$$

$$T_8(x) = 128x^8 - 256x^6 + 160x^4 - 32x^2 + 1$$

$$T_9(x) = 256x^9 - 576x^7 + 432x^5 - 120x^3 + 9x$$

$$T_{10}(x) = 512x^{10} - 1280x^8 + 1120x^6 - 400x^4 + 50x^2 - 1$$

$$T_{11}(x) = 1024x^{11} - 2816x^9 + 2816x^7 - 1232x^5 + 220x^3 - 11x$$

“BAD” POLYNOMIALS

What are those polynomials?

Chebyshev polynomials of the first kind.

$$T_0(x) = 1$$

$$T_1(x) = x$$

$$T_2(x) = 2x^2 - 1$$

$$T_3(x) = 4x^3 - 3x$$

$$T_4(x) = 8x^4 - 8x^2 + 1$$

$$T_5(x) = 16x^5 - 20x^3 + 5x$$

$$T_6(x) = 32x^6 - 48x^4 + 18x^2 - 1$$

$$T_7(x) = 64x^7 - 112x^5 + 56x^3 - 7x$$

$$T_8(x) = 128x^8 - 256x^6 + 160x^4 - 32x^2 + 1$$

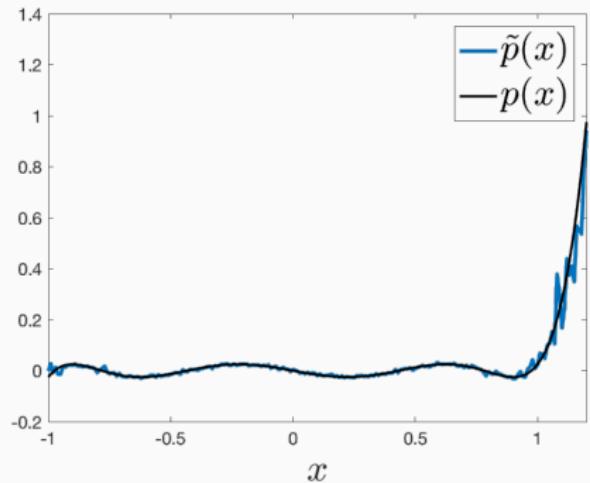
$$T_9(x) = 256x^9 - 576x^7 + 432x^5 - 120x^3 + 9x$$

$$T_{10}(x) = 512x^{10} - 1280x^8 + 1120x^6 - 400x^4 + 50x^2 - 1$$

$$T_{11}(x) = 1024x^{11} - 2816x^9 + 2816x^7 - 1232x^5 + 220x^3 - 11x$$

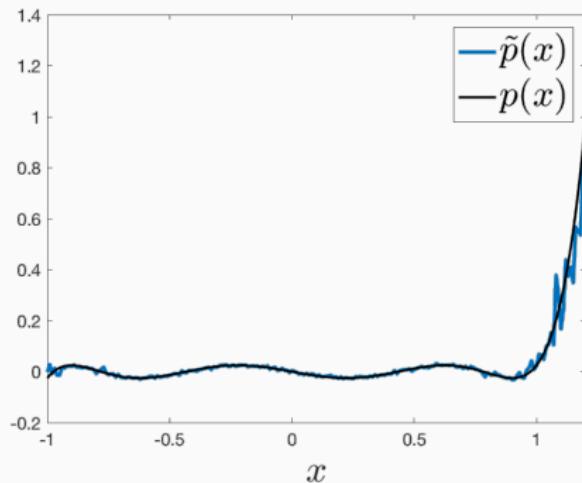
We can apply these in a stable way, using their recurrence!

“GOOD” POLYNOMIALS?



$$t_i = 2 \cdot \text{approxMult}(t_{i-1}) - t_{i-2}$$

“GOOD” POLYNOMIALS?



$$t_i = 2 \cdot \text{approxMult}(t_{i-1}) - t_{i-2}$$

Not hard to show that when computing $T_k(x)$ the error $\leq \epsilon k^2$.

KEY OBSERVATION

Chebyshev polynomials are the only hard case.

KEY OBSERVATION

Chebyshev polynomials are the only hard case.

Property: If a degree k polynomial $p(x)$ is bounded by C on $[-1, 1]$, it can be written as

$$p(x) = c_0 T_0(x) + c_1 T_1(x) + \dots + c_k T_k(x)$$

where every $|c_i| \leq 2C$.

KEY OBSERVATION

Chebyshev polynomials are the only hard case.

Property: If a degree k polynomial $p(x)$ is bounded by C on $[-1, 1]$, it can be written as

$$p(x) = c_0 T_0(x) + c_1 T_1(x) + \dots + c_k T_k(x)$$

where every $|c_i| \leq 2C$.

Total error of sum $p(x)$ is bounded by

$$2C \cdot 1^2 \epsilon + 2C \cdot 2^2 \epsilon + \dots + 2C \cdot k^2 \epsilon \leq Ck^3 \epsilon.$$

STABILITY OF LANCZOS

Same arguments extends from scalar polynomials to matrix polynomials.

STABILITY OF LANCZOS

Same arguments extends from scalar polynomials to matrix polynomials. Framework allows us to analyze Lanczos as well.

STABILITY OF LANCZOS

Same arguments extends from scalar polynomials to matrix polynomials. Framework allows us to analyze Lanczos as well.

Step 1: Lanczos stably applies Chebyshev polynomials (building on results of Paige ['71, '76, '80]).

Same arguments extends from scalar polynomials to matrix polynomials. Framework allows us to analyze Lanczos as well.

Step 1: Lanczos stably applies Chebyshev polynomials (building on results of Paige ['71, '76, '80]).

Step 2: By linearity, Lanczos stably applies polynomials bounded by C .

Same arguments extends from scalar polynomials to matrix polynomials. Framework allows us to analyze Lanczos as well.

Step 1: Lanczos stably applies Chebyshev polynomials (building on results of Paige ['71, '76, '80]).

Step 2: By linearity, Lanczos stably applies polynomials bounded by C .

Step 3: If $|f(x)| \leq C$, a good approximating polynomial has $|p(x)| \leq O(C)$, so Lanczos is stable for bounded functions.

Same arguments extends from scalar polynomials to matrix polynomials. Framework allows us to analyze Lanczos as well.

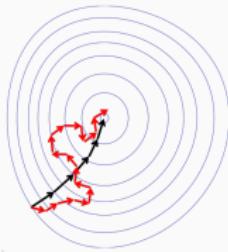
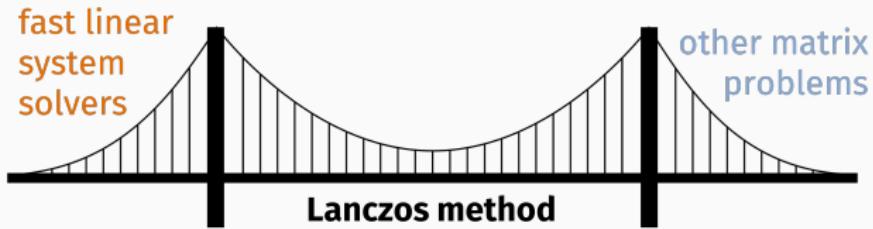
Step 1: Lanczos stably applies Chebyshev polynomials (building on results of Paige ['71, '76, '80]).

Step 2: By linearity, Lanczos stably applies polynomials bounded by C .

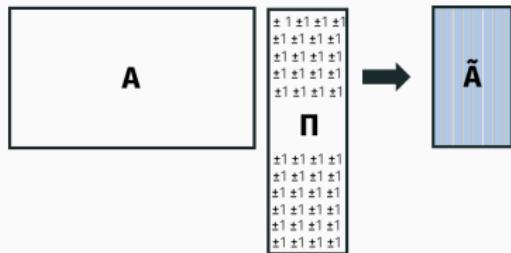
Step 3: If $|f(x)| \leq C$, a good approximating polynomial has $|p(x)| \leq O(C)$, so Lanczos is stable for bounded functions.

Use Lanczos without fear on bounded functions!

STABILITY OF LANCZOS



Stochastic Iterative
Methods



Randomized Sketching

Paper gives applications to step function, matrix exponential,
top eigenvector, etc.

Answer to old question on Lanczos in finite precision:

Theorem (Lanczos is stable for any bounded function)

If $|f(x)| \leq C$ for $x \in [\lambda_{\min}(A), \lambda_{\max}(A)]$, then if Lanczos is run for k iterations on a computer with $O(\log(nC\kappa))$ bits of precision, it outputs a vector y such that

$$\|f(A)x - y\| \leq 7k \cdot \delta_k \cdot \|x\|$$

where δ_k is the error of the best degree k uniform approximation to f .

Answer to old question on Lanczos in finite precision:

Theorem (Lanczos is stable for any bounded function)

If $|f(x)| \leq C$ for $x \in [\lambda_{\min}(A), \lambda_{\max}(A)]$, then if Lanczos is run for k iterations on a computer with $O(\log(nC\kappa))$ bits of precision, it outputs a vector y such that

$$\|f(A)x - y\| \leq 7k \cdot \delta_k \cdot \|x\|$$

where δ_k is the error of the best degree k uniform approximation to f .

- Compare to $\|f(A)x - y\| \leq 2 \cdot \delta_k \cdot \|x\|$ in exact arithmetic.

Answer to old question on Lanczos in finite precision:

Theorem (Lanczos is stable for any bounded function)

If $|f(x)| \leq C$ for $x \in [\lambda_{\min}(A), \lambda_{\max}(A)]$, then if Lanczos is run for k iterations on a computer with $O(\log(nC\kappa))$ bits of precision, it outputs a vector y such that

$$\|f(A)x - y\| \leq 7k \cdot \delta_k \cdot \|x\|$$

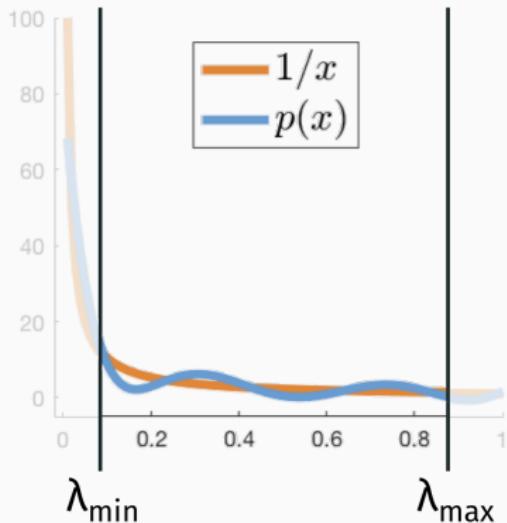
where δ_k is the error of the best degree k uniform approximation to f .

- Compare to $\|f(A)x - y\| \leq 2 \cdot \delta_k \cdot \|x\|$ in exact arithmetic.
- Matches known bound for $A^{-1}x$ (Greenbaum, '89).

NEGATIVE RESULT FOR LINEAR SYSTEMS

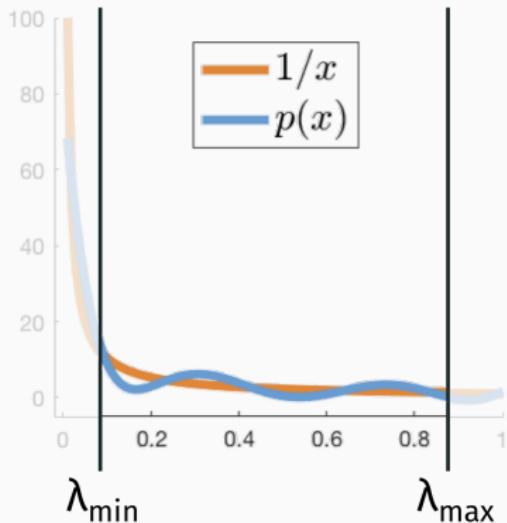
LANCZOS FOR LINEAR SYSTEMS

We proved earlier that Lanczos always matches the best uniform approximating polynomial for $f(x)$:



LANCZOS FOR LINEAR SYSTEMS

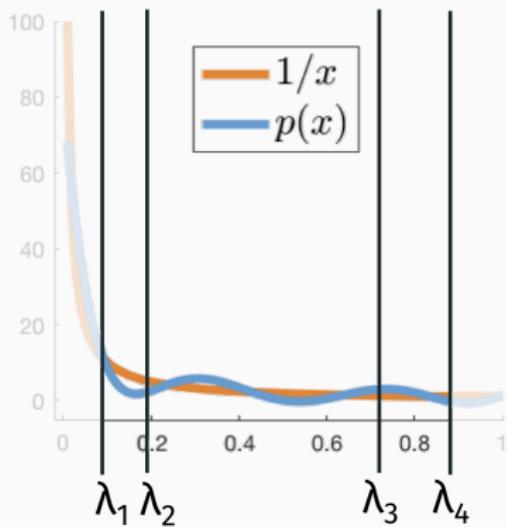
We proved earlier that Lanczos always matches the best uniform approximating polynomial for $f(x)$:



For linear systems it actually does better than that.

LANCZOS FOR LINEAR SYSTEMS

We proved earlier that Lanczos always matches the best uniform approximating polynomial for $f(x)$:



For linear systems it actually does better than that.

LANCZOS FOR LINEAR SYSTEMS

- The best **uniform** approximation to $1/x$ has degree $\sqrt{\lambda_{\max} / \lambda_{\min}} \cdot \log(1/\epsilon)$.

- The best **uniform** approximation to $1/x$ has degree $\sqrt{\lambda_{\max} / \lambda_{\min}} \cdot \log(1/\epsilon)$.
- $1/x$ can be represented exactly by a degree $n - 1$ polynomial if A only has n eigenvalues.

- The best **uniform** approximation to $1/x$ has degree $\sqrt{\lambda_{\max} / \lambda_{\min}} \cdot \log(1/\epsilon)$.
- $1/x$ can be represented exactly by a degree $n - 1$ polynomial if A only has n eigenvalues.

Claim: On exact arithmetic computers, linear systems can be solved in $O(\text{nnz}(A) \cdot n)$ time (i.e. n iterations of Lanczos)

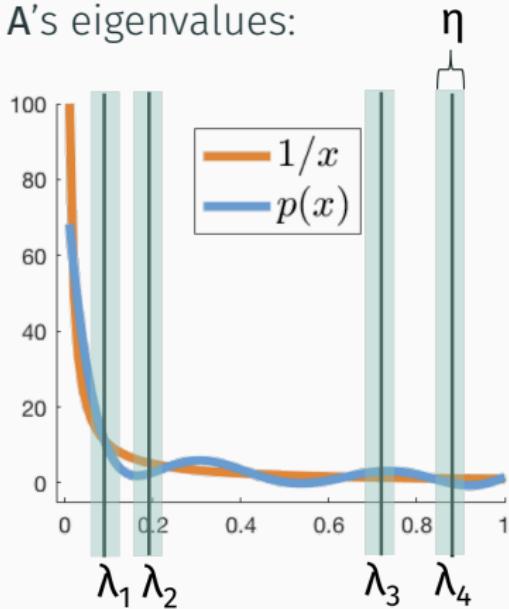
- The best **uniform** approximation to $1/x$ has degree $\sqrt{\lambda_{\max} / \lambda_{\min}} \cdot \log(1/\epsilon)$.
- $1/x$ can be represented exactly by a degree $n - 1$ polynomial if A only has n eigenvalues.

Claim: On exact arithmetic computers, linear systems can be solved in $O(\text{nnz}(A) \cdot n)$ time (i.e. n iterations of Lanczos)

Research question: To what extent does this bound hold true in finite precision? Are $n \log n$ iterations sufficient? n^2 ?

LINEAR SYSTEMS IN FINITE PRECISION

Greenbaum (1989): Finite precision Lanczos and conjugate gradient match the best polynomial approximating $1/x$ in **tiny** intervals around A's eigenvalues:



η is on the order of machine precision!

LOWER BOUND

Theorem (Stable polynomial lower bound.)

For any n , there is a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ with condition number $\lambda_{\max} / \lambda_{\min}$ such that no k degree polynomial satisfies Greenbaum's condition with error $\leq 1/3$ for all

$$k \leq (\lambda_{\max} / \lambda_{\min})^{1/5}$$

even when $\eta \leq \frac{1}{2^{n/\log \kappa}}$.

LOWER BOUND

Theorem (Stable polynomial lower bound.)

For any n , there is a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ with condition number $\lambda_{\max} / \lambda_{\min}$ such that no k degree polynomial satisfies Greenbaum's condition with error $\leq 1/3$ for all

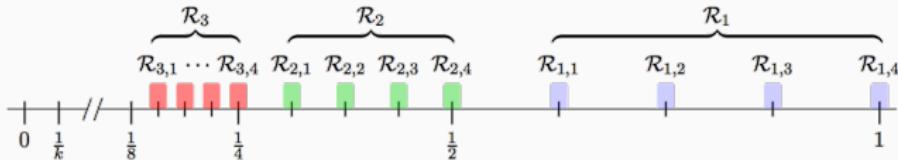
$$k \leq (\lambda_{\max} / \lambda_{\min})^{1/5}$$

even when $\eta \leq \frac{1}{2^{n/\log \kappa}}$.

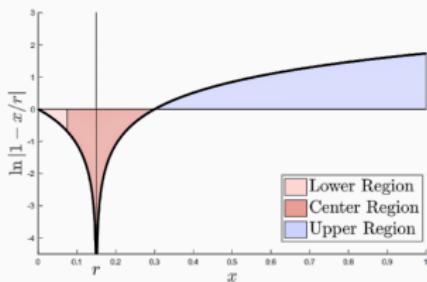
In other words, we cannot avoid polynomial dependence on condition number unless we have nearly n bits of precision.

LOWER BOUND

Construction: Eigenvalues roughly uniform on geometric scale.



Proof: Discretization of new potential function proof for Markov brother's inequality.



OPEN QUESTIONS

- Can $(\lambda_{\max} / \lambda_{\min})^{1/5}$ be tightened to $(\lambda_{\max} / \lambda_{\min})^{1/2}$
- Does Greenbaum's estimate fully characterize Lanczos?
Can the lower bound be extend to an actual runtime lower bound?
- How about for a more general class of algorithms? Any method accessing **A** only through noisy matrix-vector products?

OPEN QUESTIONS

- Can $(\lambda_{\max} / \lambda_{\min})^{1/5}$ be tightened to $(\lambda_{\max} / \lambda_{\min})^{1/2}$
- Does Greenbaum's estimate fully characterize Lanczos?
Can the lower bound be extended to an actual runtime lower bound?
- How about for a more general class of algorithms? Any method accessing \mathbf{A} only through noisy matrix-vector products?

Are $O(\text{nnz}(\mathbf{A}) \cdot n)$ time algorithms possible for linear systems?